

并行与分布式作业

“ 黄炜钊”
第六次作业

姓名：黄炜钊
班级：行政三班
学号：18340066

一、 问题描述

1、 Start from the provided skeleton code `error-test.cu` that provides some convenience macros for error checking. The macros are defined in the header file `error_checks_1.h`. Add the missing memory allocations and copies and the kernel launch and check that your code works.

1. What happens if you try to launch kernel with too large block size? When do you catch the error if you remove the `cudaDeviceSynchronize()` call?

2. What happens if you try to dereference a pointer to device memory in host code?

3. What if you try to access host memory from the kernel? Remember that you can use also `cuda-memcheck`! If you have time, you can also check what happens if you remove all error checks and do the same tests again.

2、 In this exercise we will implement a Jacobi iteration which is a very simple finite-difference scheme. Familiarize yourself with the provided skeleton. Then implement following things:

1. Write the missing CUDA kernel `sweepGPU` that implements the same algorithm as the `sweepCPU` function. Check that the reported average

difference is in the order of the numerical accuracy.

2. Experiment with different grid and block sizes and compare the execution times.

二、 解决方案

1、

(1) block size 太大会抢占计算资源, 一个报错例子如下:

```
Error: vector_add kernel at 0_4323.cu(86): invalid configuration argument  
yhrun: error: gn07: task 1: Exited with exit code 1
```

删除了 cudaDeviceSynchronize() call 之后, 报错信息如下:

```
Error: vector_add kernel at 0_4323.cu(86): invalid configuration argument  
yhrun: error: gn07: task 0: Exited with exit code 1
```

(2) 当运行到解除引用的时候会发生报错, 具体信息如下:

```
Error: vector_add kernel at 0_4323.cu(85): invalid configuration argument  
yhrun: error: gn07: task 0: Exited with exit code 1
```

(3) 在编译的时候就会报错, 具体信息如下:

```
0_4323.cu(48): error: identifier "hb" is undefined in device code  
  
1 error detected in the compilation of "/tmp/tmpxft_0000440e_00000000-11_0_4323.cpp2.i".  
  
slurmd[gn07]: execve(): 0_4323.cu.out: No such file or directory  
yhrun: error: gn07: task 0: Exited with exit code 2  
rm: cannot remove '0_4323.cu.out': No such file or directory
```

代码如下:

```
#include <stdio.h>  
#include <math.h>  
// #include "error_checks.h" // Macros CUDA_CHECK and CHECK_ERROR_MSG  
// This header provides two helper macros for error checking  
// See the exercise skeletons and answers for usage examples.  
  
#ifndef COURSE_UTIL_H_
```

```

#define COURSE_UTIL_H_

#include <stdio.h>
#include <stdlib.h>

#define CUDA_CHECK(errang) __checkErrorFunc(errang, __FILE__, __LINE__)
#define CHECK_ERROR_MSG(errstr) __checkErrMsgFunc(errstr, __FILE__, __LINE__)

inline void __checkErrorFunc(cudaError_t errang, const char *file,
                             const int line)
{
    if (errang)
    {
        fprintf(stderr, "Error at %s(%i)\n", file, line);
        exit(EXIT_FAILURE);
    }
}

inline void __checkErrMsgFunc(const char *errstr, const char *file,
                              const int line)
{
    cudaError_t err = cudaGetLastError();
    if (err != cudaSuccess)
    {
        fprintf(stderr, "Error: %s at %s(%i): %s\n",
                errstr, file, line, cudaGetErrorString(err));
        exit(EXIT_FAILURE);
    }
}

#endif

__global__ void vector_add(double *C, const double *A, const double *B,
                           int N)
{
    // Add the kernel code
    int idx = blockIdx.x * blockDim.x + threadIdx.x;

    // Do not try to access past the allocated memory
    if (idx < N)
    {
        C[idx] = A[idx] + B[idx];
    }
}

```

```

}

int main(void)
{
    const int N = 20;
    const int ThreadsInBlock = 128;
    double *dA, *dB, *dC;
    double hA[N], hB[N], hC[N];

    for (int i = 0; i < N; ++i)
    {
        hA[i] = (double)i;
        hB[i] = (double)i * i;
    }

    /*
        Add memory allocations and copies. Wrap your runtime function
        calls with CUDA_CHECK( ) macro
    */
    CUDA_CHECK(cudaMalloc((void **)&dA, sizeof(double) * N));
    CUDA_CHECK(cudaMalloc((void **)&dB, sizeof(double) * N));
    CUDA_CHECK(cudaMalloc((void **)&dC, sizeof(double) * N));
    CUDA_CHECK(cudaMemcpy(dA, hA, sizeof(double) * N, cudaMemcpyHostToDevice));
    CUDA_CHECK(cudaMemcpy(dB, hB, sizeof(double) * N, cudaMemcpyHostToDevice));
    /*#error Add the remaining memory allocations and copies

    // Note the maximum size of threads in a block
    dim3 grid, threads;

    //// Add the kernel call here
    vector_add<<<1, 32>>>>(dC, dA, dB, N);
    /*#error Add the CUDA kernel call

    // Here we add an explicit synchronization so that we catch errors
    // as early as possible. Don't do this in production code!
    cudaDeviceSynchronize();
    CHECK_ERROR_MSG("vector_add kernel");

    //// Copy back the results and free the device memory
    CUDA_CHECK(cudaMemcpy(hC, dC, sizeof(double) * N, cudaMemcpyDeviceToHost));
    CUDA_CHECK(cudaFree(dA));

```

```

    CUDA_CHECK(cudaFree(dB));
    CUDA_CHECK(cudaFree(dC));
    //error Copy back the results and free the allocated memory

    for (int i = 0; i < N; i++)
        printf("%5.1f\n", hC[i]);

    return 0;
}

```

2、代码如下：

```

//err_checker.h
// This header provides two helper macros for error checking
// See the exercise skeletons and answers for usage examples.

#ifndef COURSE_UTIL_H_
#define COURSE_UTIL_H_

#include <stdio.h>
#include <stdlib.h>

#define CUDA_CHECK(errarg) __checkErrorFunc(errarg, __FILE__, __LINE__)
#define CHECK_ERROR_MSG(errstr) __checkErrMsgFunc(errstr, __FILE__, __LINE__)

inline void __checkErrorFunc(cudaError_t errarg, const char *file,
                             const int line)
{
    if (errarg)
    {
        fprintf(stderr, "Error at %s(%i)\n", file, line);
        exit(EXIT_FAILURE);
    }
}

inline void __checkErrMsgFunc(const char *errstr, const char *file,
                              const int line)
{
    cudaError_t err = cudaGetLastError();
    if (err != cudaSuccess)
    {
        fprintf(stderr, "Error: %s at %s(%i): %s\n",
                errstr, file, line, cudaGetErrorString(err));
    }
}

```

```

        exit(EXIT_FAILURE);
    }
}

#endif
//jacbi.h
#ifndef EX3_H_
#define EX3_H_

#include <thrust/device_vector.h>
#include <thrust/functional.h>
#include <thrust/transform_reduce.h>
#include <thrust/iterator/zip_iterator.h>

// Helper function prototypes
double compareArrays(const double *a, const double *b, int N);
double diffCPU(const double *a, const double *b, int N);
void sweepCPU(double *phi, const double *phiPrev,
              const double *source, double h2, int N);

/* -----
-----
    EXTRACURRICULAR ACTIVITIES

    This part provides the reduction operation (in this case summation o
f
    difference of two arrays) using thrust library. Thrust mimics the
    syntax and design of standard template library (STL) of C++. Thrust
is
    also a part of CUDA 4 SDK.
    More information can be found from thrust home page
    http://code.google.com/p/thrust/
    -----
--- */

template <typename T>
class square_diff_thr : public thrust::unary_function<thrust::tuple<T,
T>, T>
{
public:
    __host__ __device__
    T
    operator()(const thrust::tuple<T, T> &x) const
    {

```

```

        return (thrust::get<1>(x) - thrust::get<0>(x)) *
               (thrust::get<1>(x) - thrust::get<0>(x));
    }
};

template <typename T>
class square_thr : public thrust::unary_function<T, T>
{
public:
    __host__ __device__
    T
    operator()(const T &x) const
    {
        return x * x;
    }
};

template <typename T>
T diffGPU(T *A_d, T *B_d, int N)
{
    typedef thrust::device_ptr<T> FloatIterator;
    typedef thrust::tuple<FloatIterator, FloatIterator> IteratorTuple;
    typedef thrust::zip_iterator<IteratorTuple> ZipIterator;

    thrust::device_ptr<T> A_ptr(A_d);
    thrust::device_ptr<T> B_ptr(B_d);

    ZipIterator first =
        thrust::make_zip_iterator(thrust::make_tuple(A_ptr, B_ptr));
    ZipIterator last =
        thrust::make_zip_iterator(thrust::make_tuple(A_ptr + N * N,
                                                       B_ptr + N * N));

    T a1 = thrust::transform_reduce(first, last, square_diff_thr<T>(),
                                     static_cast<T>(0), thrust::plus<T>());
    T a2 = thrust::transform_reduce(B_ptr, B_ptr + N * N,
                                     square_thr<T>(), static_cast<T>(0),
                                     thrust::plus<T>());

    return sqrt(a1 / a2);
}

#endif // EX3_H_

```



```

//jacobi.cu
#include <time.h>
#include <stdio.h>
//#include "jacobi.h"
//#include "error_checks.h"

// Change this to 0 if CPU reference result is not needed
#define COMPUTE_CPU_REFERENCE 1
#define MAX_ITERATIONS 3000

// CPU kernel
void sweepCPU(double *phi, const double *phiPrev, const double *source,
              double h2, int N)
{
    int i, j;
    int index, i1, i2, i3, i4;

    for (j = 1; j < N - 1; j++)
    {
        for (i = 1; i < N - 1; i++)
        {
            index = i + j * N;
            i1 = (i - 1) + j * N;
            i2 = (i + 1) + j * N;
            i3 = i + (j - 1) * N;
            i4 = i + (j + 1) * N;
            phi[index] = 0.25 * (phiPrev[i1] + phiPrev[i2] +
                                phiPrev[i3] + phiPrev[i4] -
                                h2 * source[index]);
        }
    }
}

// GPU kernel
__global__ void sweepGPU(double *phi, const double *phiPrev, const double *source, double h2, int N)
{
    // #error Add here the GPU version of the update routine (see sweep
    CPU above)
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    if (i > 0 && j > 0 && i < N - 1 && j < N - 1)
    { // be careful!

```

```

        int index = i + j * N;
        int i1 = (i - 1) + j * N;
        int i2 = (i + 1) + j * N;
        int i3 = i + (j - 1) * N;
        int i4 = i + (j + 1) * N;
        phi[index] = 0.25 * (phiPrev[i1] + phiPrev[i2] + phiPrev[i3] +
phiPrev[i4] - h2 * source[index]);
    }
}

double compareArrays(const double *a, const double *b, int N)
{
    double error = 0.0;
    int i;
    for (i = 0; i < N * N; i++)
    {
        error += fabs(a[i] - b[i]);
    }
    return error / (N * N);
}

double diffCPU(const double *phi, const double *phiPrev, int N)
{
    int i;
    double sum = 0;
    double diffsum = 0;

    for (i = 0; i < N * N; i++)
    {
        diffsum += (phi[i] - phiPrev[i]) * (phi[i] - phiPrev[i]);
        sum += phi[i] * phi[i];
    }

    return sqrt(diffsum / sum);
}

int main()
{
    clock_t t1, t2; // Structs for timing
    const int N = 512;
    double h = 1.0 / (N - 1);
    int iterations;
    const double tolerance = 5e-4; // Stopping condition
    int i, j, index;

```

```

const int blocksize = 16;

double *phi = new double[N * N];
double *phiPrev = new double[N * N];
double *source = new double[N * N];
double *phi_cuda = new double[N * N];

double *phi_d, *phiPrev_d, *source_d;
// Size of the arrays in bytes
const int size = N * N * sizeof(double);
double diff;

// Source initialization
for (i = 0; i < N; i++)
{
    for (j = 0; j < N; j++)
    {
        double x, y;
        x = (i - N / 2) * h;
        y = (j - N / 2) * h;
        index = j + i * N;
        if (((x - 0.25) * (x - 0.25) + y * y) < 0.1 * 0.1)
            source[index] = 1e10 * h * h;
        else if (((x + 0.25) * (x + 0.25) + y * y) < 0.1 * 0.1)
            source[index] = -1e10 * h * h;
        else
            source[index] = 0.0;
    }
}

CUDA_CHECK(cudaMalloc((void **)&source_d, size));
CUDA_CHECK(cudaMemcpy(source_d, source, size, cudaMemcpyHostToDevice));

// Reset values to zero
for (i = 0; i < N; i++)
{
    for (j = 0; j < N; j++)
    {
        index = j + i * N;
        phi[index] = 0.0;
        phiPrev[index] = 0.0;
    }
}

```

```

    }

    CUDA_CHECK(cudaMalloc((void **)&phi_d, size));
    CUDA_CHECK(cudaMalloc((void **)&phiPrev_d, size));
    CUDA_CHECK(cudaMemcpy(phi_d, phi, size, cudaMemcpyHostToDevice));
    CUDA_CHECK(cudaMemcpy(phiPrev_d, phiPrev, size, cudaMemcpyHostToDevice));

    // CPU version
    if (COMPUTE_CPU_REFERENCE)
    {
        t1 = clock();

        // Do sweeps untill difference is under the tolerance
        diff = tolerance * 2;
        iterations = 0;
        while (diff > tolerance && iterations < MAX_ITERATIONS)
        {
            sweepCPU(phiPrev, phi, source, h * h, N);
            sweepCPU(phi, phiPrev, source, h * h, N);

            iterations += 2;
            if (iterations % 100 == 0)
            {
                diff = diffCPU(phi, phiPrev, N);
                printf("%d %g\n", iterations, diff);
            }
        }
        t2 = clock();
        printf("CPU Jacobi: %g ms, %d iterations\n",
            t2 - t1,
            iterations);
    }

    // GPU version

    dim3 dimBlock(blocksize, blocksize);
    dim3 dimGrid((N + blocksize - 1) / blocksize, (N + blocksize - 1) /
        blocksize);

    //do sweeps until diff under tolerance
    diff = tolerance * 2;
    iterations = 0;

```

```

t1 = clock();

while (diff > tolerance && iterations < MAX_ITERATIONS)
{
    // See above how the CPU update kernel is called
    // and implement similar calling sequence for the GPU code

    ///// Add routines here
    sweepGPU<<<dimGrid, dimBlock>>>(phiPrev_d, phi_d, source, h * h
, N);
    sweepGPU<<<dimGrid, dimBlock>>>(phi_d, phiPrev_d, source, h * h
, N);
    ///#error Add GPU kernel calls here (see CPU version above)

    iterations += 2;

    if (iterations % 100 == 0)
    {
        // diffGPU is defined in the header file, it uses
        // Thrust library for reduction computation
        diff = diffGPU<double>(phiPrev_d, phi_d, N);
        CHECK_ERROR_MSG("Difference computation");
        printf("%d %g\n", iterations, diff);
    }
}

///// Add here the routine to copy back the results
CUDA_CHECK(cudaMemcpy(phi, phi_d, size, cudaMemcpyDeviceToHost));
CUDA_CHECK(cudaMemcpy(phiPrev, phiPrev_d, size, cudaMemcpyDeviceToH
ost));
///#error Copy back the results

t2 = clock();
printf("GPU Jacobi: %g ms, %d iterations\n",
      t2 - t1,
      iterations);

///// Add here the clean up code for all allocated CUDA resources
CUDA_CHECK(cudaFree(phi_d));
CUDA_CHECK(cudaFree(phiPrev_d));
CUDA_CHECK(cudaFree(source_d));
///#error Add here the clean up code

if (COMPUTE_CPU_REFERENCE)

```

```
{  
    printf("Average difference is %g\n", compareArrays(phi, phi_cuda,  
a, N));  
}  
  
delete[] phi;  
delete[] phi_cuda;  
delete[] phiPrev;  
delete[] source;  
  
return EXIT_SUCCESS;  
}
```